Adapters Learn Less and Forget Less

Moreau, Patel et al. · 2024

Fictional sample paper generated for the ClaimLens demo.

Results for individual tasks appear in the appendix. We release the evaluation scripts so that every number can be reproduced. Training code is released with the paper. We release the evaluation scripts so that every number can be reproduced. We discuss the limitations of this protocol in the final section.

All runs share the same learning-rate schedule with a short linear warmup. We evaluate the final checkpoint rather than the best one seen during training. We report wall-clock time on a single node of eight accelerators. We follow the standard training recipe and report the mean over three seeds.

The comparison holds the data and the evaluation harness constant. This setting matches the one most often used in practice. Unless noted otherwise, hyperparameters are tuned on a held-out validation split. Variance across seeds is small compared with the differences we report.

Ablations in the appendix vary one factor at a time. Results for individual tasks appear in the appendix. We release the evaluation scripts so that every number can be reproduced. We discuss the limitations of this protocol in the final section.

We evaluate the final checkpoint rather than the best one seen during training. We report wall-clock time on a single node of eight accelerators. We follow the standard training recipe and report the mean over three seeds. Each configuration is trained for the same number of optimizer steps.

This setting matches the one most often used in practice. Unless noted otherwise, hyperparameters are tuned on a held-out validation split. Variance across seeds is small compared with the differences we report. Prior work reached similar conclusions on smaller models.

The comparison holds the data and the evaluation harness constant. This setting matches the one most often used in practice. Unless noted otherwise, hyperparameters are tuned on a held-out validation split. Variance across seeds is small compared with the differences we report.

Ablations in the appendix vary one factor at a time. Results for individual tasks appear in the appendix. We release the evaluation scripts so that every number can be reproduced. We discuss the limitations of this protocol in the final section.

We evaluate the final checkpoint rather than the best one seen during training. We report wall-clock time on a single node of eight accelerators. We follow the standard training recipe and report the mean over three seeds. Each configuration is trained for the same number of optimizer steps.

This setting matches the one most often used in practice. Unless noted otherwise, hyperparameters are tuned on a held-out validation split. Variance across seeds is small compared with the differences we report. Prior work reached similar conclusions on smaller models.

Results for individual tasks appear in the appendix. We release the evaluation scripts so that every number can be reproduced. We discuss the limitations of this protocol in the final section. We keep the tokenizer and context length fixed across every configuration.

We report wall-clock time on a single node of eight accelerators. We follow the standard training recipe and report the mean over three seeds. Each configuration is trained for the same number of optimizer steps. Training uses mixed precision with gradient clipping at a norm of one.

Unless noted otherwise, hyperparameters are tuned on a held-out validation split. Variance across seeds is small compared with the differences we report. We compare LoRA against full fine-tuning on the same data. Variance across seeds is small compared with the differences we report. Prior work reached similar conclusions on smaller models.

Ablations in the appendix vary one factor at a time. Results for individual tasks appear in the appendix. We release the evaluation scripts so that every number can be reproduced. We discuss the limitations of this protocol in the final section.

We evaluate the final checkpoint rather than the best one seen during training. We report wall-clock time on a single node of eight accelerators. We follow the standard training recipe and report the mean over three seeds. Each configuration is trained for the same number of optimizer steps.

This setting matches the one most often used in practice. Unless noted otherwise, hyperparameters are tuned on a held-out validation split. Variance across seeds is small compared with the differences we report. Prior work reached similar conclusions on smaller models.

Results for individual tasks appear in the appendix. We release the evaluation scripts so that every number can be reproduced. We discuss the limitations of this protocol in the final section. We keep the tokenizer and context length fixed across every configuration.

We report wall-clock time on a single node of eight accelerators. We follow the standard training recipe and report the mean over three seeds. Each configuration is trained for the same number of optimizer steps. Training uses mixed precision with gradient clipping at a norm of one.

We release the evaluation scripts so that every number can be reproduced. We discuss the limitations of this protocol in the final section. All models start from the same 7B base model. We discuss the limitations of this protocol in the final section. We keep the tokenizer and context length fixed across every configuration.

We evaluate the final checkpoint rather than the best one seen during training. We report wall-clock time on a single node of eight accelerators. We follow the standard training recipe and report the mean over three seeds. Each configuration is trained for the same number of optimizer steps.

This setting matches the one most often used in practice. Unless noted otherwise, hyperparameters are tuned on a held-out validation split. Variance across seeds is small compared with the differences we report. Prior work reached similar conclusions on smaller models.

Results for individual tasks appear in the appendix. We release the evaluation scripts so that every number can be reproduced. We discuss the limitations of this protocol in the final section. We keep the tokenizer and context length fixed across every configuration.

We report wall-clock time on a single node of eight accelerators. We follow the standard training recipe and report the mean over three seeds. Each configuration is trained for the same number of optimizer steps. Training uses mixed precision with gradient clipping at a norm of one.

Unless noted otherwise, hyperparameters are tuned on a held-out validation split. Variance across seeds is small compared with the differences we report. Prior work reached similar conclusions on smaller models. All runs share the same learning-rate schedule with a short linear warmup.

We follow the standard training recipe and report the mean over three seeds. Each configuration is trained for the same number of optimizer steps. We report pass@1 for code tasks and accuracy for math. Each configuration is trained for the same number of optimizer steps. Training uses mixed precision with gradient clipping at a norm of one.

Variance across seeds is small compared with the differences we report. Prior work reached similar conclusions on smaller models. Figure 2 covers the code and math benchmarks. Prior work reached similar conclusions on smaller models. All runs share the same learning-rate schedule with a short linear warmup.

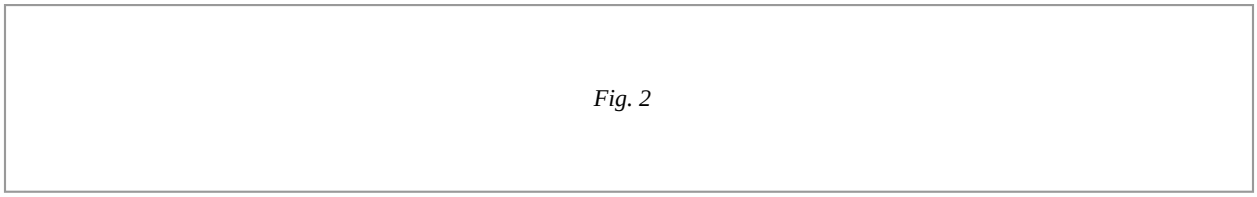

We discuss the limitations of this protocol in the final section. We keep the tokenizer and context length fixed across every configuration. On code and math, LoRA substantially underperforms full fine-tuning even with long training runs. We keep the tokenizer and context length fixed across every configuration. The comparison holds the data and the evaluation harness constant.

Each configuration is trained for the same number of optimizer steps. Training uses mixed precision with gradient clipping at a norm of one. On code generation, LoRA substantially underperforms full fine-tuning. Training uses mixed precision with gradient clipping at a norm of one. Ablations in the appendix vary one factor at a time.

This setting matches the one most often used in practice. Unless noted otherwise, hyperparameters are tuned on a held-out validation split. Variance across seeds is small compared with the differences we report. Prior work reached similar conclusions on smaller models.

Results for individual tasks appear in the appendix. We release the evaluation scripts so that every number can be reproduced. We discuss the limitations of this protocol in the final section. We keep the tokenizer and context length fixed across every configuration.

Results for individual tasks appear in the appendix. We release the evaluation scripts so that every number can be reproduced. We discuss the limitations of this protocol in the final section. We keep the tokenizer and context length fixed across every configuration.

We report wall-clock time on a single node of eight accelerators. We follow the standard training recipe and report the mean over three seeds. Each configuration is trained for the same number of optimizer steps. Training uses mixed precision with gradient clipping at a norm of one.

Unless noted otherwise, hyperparameters are tuned on a held-out validation split. Variance across seeds is small compared with the differences we report. Prior work reached similar conclusions on smaller models. All runs share the same learning-rate schedule with a short linear warmup.

We release the evaluation scripts so that every number can be reproduced. We discuss the limitations of this protocol in the final section. We keep the tokenizer and context length fixed across every configuration. The comparison holds the data and the evaluation harness constant.

We follow the standard training recipe and report the mean over three seeds. Each configuration is trained for the same number of optimizer steps. Training uses mixed precision with gradient clipping at a norm of one. Ablations in the appendix vary one factor at a time.

Variance across seeds is small compared with the differences we report. Prior work reached similar conclusions on smaller models. All runs share the same learning-rate schedule with a short linear warmup. We evaluate the final checkpoint rather than the best one seen during training.

We report wall-clock time on a single node of eight accelerators. We follow the standard training recipe and report the mean over three seeds. Each configuration is trained for the same number of optimizer steps. Training uses mixed precision with gradient clipping at a norm of one.

Unless noted otherwise, hyperparameters are tuned on a held-out validation split. Variance across seeds is small compared with the differences we report. Prior work reached similar conclusions on smaller models. All runs share the same learning-rate schedule with a short linear warmup.

We release the evaluation scripts so that every number can be reproduced. We discuss the limitations of this protocol in the final section. We keep the tokenizer and context length fixed across every configuration. The comparison holds the data and the evaluation harness constant.

We follow the standard training recipe and report the mean over three seeds. Each configuration is trained for the same number of optimizer steps. Training uses mixed precision with gradient clipping at a norm of one. Ablations in the appendix vary one factor at a time.

Variance across seeds is small compared with the differences we report. Prior work reached similar conclusions on smaller models. All runs share the same learning-rate schedule with a short linear warmup. We evaluate the final checkpoint rather than the best one seen during training.

We discuss the limitations of this protocol in the final section. We keep the tokenizer and context length fixed across every configuration. The comparison holds the data and the evaluation harness constant. This setting matches the one most often used in practice.

Each configuration is trained for the same number of optimizer steps. Training uses mixed precision with gradient clipping at a norm of one. LoRA better preserves performance on tasks outside the fine-tuning domain. Training uses mixed precision with gradient clipping at a norm of one. Ablations in the appendix vary one factor at a time.

Unless noted otherwise, hyperparameters are tuned on a held-out validation split. Variance across seeds is small compared with the differences we report. Prior work reached similar conclusions on smaller models. All runs share the same learning-rate schedule with a short linear warmup.

We release the evaluation scripts so that every number can be reproduced. We discuss the limitations of this protocol in the final section. We keep the tokenizer and context length fixed across every configuration. The comparison holds the data and the evaluation harness constant.

We follow the standard training recipe and report the mean over three seeds. Each configuration is trained for the same number of optimizer steps. Training uses mixed precision with gradient clipping at a norm of one. Ablations in the appendix vary one factor at a time.

Variance across seeds is small compared with the differences we report. Prior work reached similar conclusions on smaller models. All runs share the same learning-rate schedule with a short linear warmup. We evaluate the final checkpoint rather than the best one seen during training.

We discuss the limitations of this protocol in the final section. We keep the tokenizer and context length fixed across every configuration. The comparison holds the data and the evaluation harness constant. This setting matches the one most often used in practice.

We release the evaluation scripts so that every number can be reproduced. We discuss the limitations of this protocol in the final section. We keep the tokenizer and context length fixed across every configuration. The comparison holds the data and the evaluation harness constant.

We follow the standard training recipe and report the mean over three seeds. Each configuration is trained for the same number of optimizer steps. Training uses mixed precision with gradient clipping at a norm of one. Ablations in the appendix vary one factor at a time.

Variance across seeds is small compared with the differences we report. Prior work reached similar conclusions on smaller models. All runs share the same learning-rate schedule with a short linear warmup. We evaluate the final checkpoint rather than the best one seen during training.

We discuss the limitations of this protocol in the final section. We keep the tokenizer and context length fixed across every configuration. The comparison holds the data and the evaluation harness constant. This setting matches the one most often used in practice.

Each configuration is trained for the same number of optimizer steps. Training uses mixed precision with gradient clipping at a norm of one. Ablations in the appendix vary one factor at a time. Results for individual tasks appear in the appendix.

Prior work reached similar conclusions on smaller models. All runs share the same learning-rate schedule with a short linear warmup. We evaluate the final checkpoint rather than the best one seen during training. We report wall-clock time on a single node of eight accelerators.

We follow the standard training recipe and report the mean over three seeds. Each configuration is trained for the same number of optimizer steps. Training uses mixed precision with gradient clipping at a norm of one. Ablations in the appendix vary one factor at a time.

Variance across seeds is small compared with the differences we report. Prior work reached similar conclusions on smaller models. All runs share the same learning-rate schedule with a short linear warmup. We evaluate the final checkpoint rather than the best one seen during training.

We discuss the limitations of this protocol in the final section. We keep the tokenizer and context length fixed across every configuration. The comparison holds the data and the evaluation harness constant. This setting matches the one most often used in practice.

Each configuration is trained for the same number of optimizer steps. Training uses mixed precision with gradient clipping at a norm of one. Ablations in the appendix vary one factor at a time. Results for individual tasks appear in the appendix.

Prior work reached similar conclusions on smaller models. All runs share the same learning-rate schedule with a short linear warmup. We evaluate the final checkpoint rather than the best one seen during training. We report wall-clock time on a single node of eight accelerators.

We keep the tokenizer and context length fixed across every configuration. The comparison holds the data and the evaluation harness constant. This setting matches the one most often used in practice. Unless noted otherwise, hyperparameters are tuned on a held-out validation split.

Variance across seeds is small compared with the differences we report. Prior work reached similar conclusions on smaller models. All runs share the same learning-rate schedule with a short linear warmup. We evaluate the final checkpoint rather than the best one seen during training.

We discuss the limitations of this protocol in the final section. We keep the tokenizer and context length fixed across every configuration. The comparison holds the data and the evaluation harness constant. This setting matches the one most often used in practice.

Each configuration is trained for the same number of optimizer steps. Training uses mixed precision with gradient clipping at a norm of one. Ablations in the appendix vary one factor at a time. Results for individual tasks appear in the appendix.

Prior work reached similar conclusions on smaller models. All runs share the same learning-rate schedule with a short linear warmup. We evaluate the final checkpoint rather than the best one seen during training. We report wall-clock time on a single node of eight accelerators.

We keep the tokenizer and context length fixed across every configuration. The comparison holds the data and the evaluation harness constant. This setting matches the one most often used in practice. Unless noted otherwise, hyperparameters are tuned on a held-out validation split.

Training uses mixed precision with gradient clipping at a norm of one. Ablations in the appendix vary one factor at a time. Results for individual tasks appear in the appendix. We release the evaluation scripts so that every number can be reproduced.

We discuss the limitations of this protocol in the final section. We keep the tokenizer and context length fixed across every configuration. The comparison holds the data and the evaluation harness constant. This setting matches the one most often used in practice.

Each configuration is trained for the same number of optimizer steps. Training uses mixed precision with gradient clipping at a norm of one. Ablations in the appendix vary one factor at a time. Results for individual tasks appear in the appendix.

Prior work reached similar conclusions on smaller models. All runs share the same learning-rate schedule with a short linear warmup. We evaluate the final checkpoint rather than the best one seen during training. We report wall-clock time on a single node of eight accelerators.

We keep the tokenizer and context length fixed across every configuration. The comparison holds the data and the evaluation harness constant. This setting matches the one most often used in practice. Unless noted otherwise, hyperparameters are tuned on a held-out validation split.

Training uses mixed precision with gradient clipping at a norm of one. Ablations in the appendix vary one factor at a time. Results for individual tasks appear in the appendix. We release the evaluation scripts so that every number can be reproduced.

All runs share the same learning-rate schedule with a short linear warmup. We evaluate the final checkpoint rather than the best one seen during training. We report wall-clock time on a single node of eight accelerators. We follow the standard training recipe and report the mean over three seeds.

Each configuration is trained for the same number of optimizer steps. Training uses mixed precision with gradient clipping at a norm of one. Ablations in the appendix vary one factor at a time. Results for individual tasks appear in the appendix.

Prior work reached similar conclusions on smaller models. All runs share the same learning-rate schedule with a short linear warmup. We evaluate the final checkpoint rather than the best one seen during training. We report wall-clock time on a single node of eight accelerators.

We keep the tokenizer and context length fixed across every configuration. The comparison holds the data and the evaluation harness constant. This setting matches the one most often used in practice. Unless noted otherwise, hyperparameters are tuned on a held-out validation split.

Training uses mixed precision with gradient clipping at a norm of one. Ablations in the appendix vary one factor at a time. Results for individual tasks appear in the appendix. We release the evaluation scripts so that every number can be reproduced.

All runs share the same learning-rate schedule with a short linear warmup. We evaluate the final checkpoint rather than the best one seen during training. We report wall-clock time on a single node of eight accelerators. We follow the standard training recipe and report the mean over three seeds.

The comparison holds the data and the evaluation harness constant. This setting matches the one most often used in practice. Unless noted otherwise, hyperparameters are tuned on a held-out validation split. Variance across seeds is small compared with the differences we report.

Prior work reached similar conclusions on smaller models. All runs share the same learning-rate schedule with a short linear warmup. We evaluate the final checkpoint rather than the best one seen during training. We report wall-clock time on a single node of eight accelerators.

We keep the tokenizer and context length fixed across every configuration. The comparison holds the data and the evaluation harness constant. This setting matches the one most often used in practice. Unless noted otherwise, hyperparameters are tuned on a held-out validation split.

Training uses mixed precision with gradient clipping at a norm of one. Ablations in the appendix vary one factor at a time. Results for individual tasks appear in the appendix. We release the evaluation scripts so that every number can be reproduced.

All runs share the same learning-rate schedule with a short linear warmup. We evaluate the final checkpoint rather than the best one seen during training. We report wall-clock time on a single node of eight accelerators. We follow the standard training recipe and report the mean over three seeds.

The comparison holds the data and the evaluation harness constant. This setting matches the one most often used in practice. Unless noted otherwise, hyperparameters are tuned on a held-out validation split. Variance across seeds is small compared with the differences we report.

Ablations in the appendix vary one factor at a time. Results for individual tasks appear in the appendix. We release the evaluation scripts so that every number can be reproduced. We discuss the limitations of this protocol in the final section.

We keep the tokenizer and context length fixed across every configuration. The comparison holds the data and the evaluation harness constant. This setting matches the one most often used in practice. Unless noted otherwise, hyperparameters are tuned on a held-out validation split.

Training uses mixed precision with gradient clipping at a norm of one. Ablations in the appendix vary one factor at a time. Results for individual tasks appear in the appendix. We release the evaluation scripts so that every number can be reproduced.

All runs share the same learning-rate schedule with a short linear warmup. We evaluate the final checkpoint rather than the best one seen during training. We report wall-clock time on a single node of eight accelerators. We follow the standard training recipe and report the mean over three seeds.

The comparison holds the data and the evaluation harness constant. This setting matches the one most often used in practice. Unless noted otherwise, hyperparameters are tuned on a held-out validation split. Variance across seeds is small compared with the differences we report.

Ablations in the appendix vary one factor at a time. Results for individual tasks appear in the appendix. We release the evaluation scripts so that every number can be reproduced. We discuss the limitations of this protocol in the final section.

We evaluate the final checkpoint rather than the best one seen during training. We report wall-clock time on a single node of eight accelerators. We follow the standard training recipe and report the mean over three seeds. Each configuration is trained for the same number of optimizer steps.

Training uses mixed precision with gradient clipping at a norm of one. Ablations in the appendix vary one factor at a time. Results for individual tasks appear in the appendix. We release the evaluation scripts so that every number can be reproduced.

All runs share the same learning-rate schedule with a short linear warmup. We evaluate the final checkpoint rather than the best one seen during training. We report wall-clock time on a single node of eight accelerators. We follow the standard training recipe and report the mean over three seeds.

The comparison holds the data and the evaluation harness constant. This setting matches the one most often used in practice. Unless noted otherwise, hyperparameters are tuned on a held-out validation split. Variance across seeds is small compared with the differences we report.

Ablations in the appendix vary one factor at a time. Results for individual tasks appear in the appendix. We release the evaluation scripts so that every number can be reproduced. We discuss the limitations of this protocol in the final section.

We evaluate the final checkpoint rather than the best one seen during training. We report wall-clock time on a single node of eight accelerators. We follow the standard training recipe and report the mean over three seeds. Each configuration is trained for the same number of optimizer steps.

This setting matches the one most often used in practice. Unless noted otherwise, hyperparameters are tuned on a held-out validation split. Variance across seeds is small compared with the differences we report. Prior work reached similar conclusions on smaller models.

All runs share the same learning-rate schedule with a short linear warmup. We evaluate the final checkpoint rather than the best one seen during training. We report wall-clock time on a single node of eight accelerators. We follow the standard training recipe and report the mean over three seeds.

The comparison holds the data and the evaluation harness constant. This setting matches the one most often used in practice. Unless noted otherwise, hyperparameters are tuned on a held-out validation split. Variance across seeds is small compared with the differences we report.

Ablations in the appendix vary one factor at a time. Results for individual tasks appear in the appendix. We release the evaluation scripts so that every number can be reproduced. We discuss the limitations of this protocol in the final section.

We evaluate the final checkpoint rather than the best one seen during training. We report wall-clock time on a single node of eight accelerators. We follow the standard training recipe and report the mean over three seeds. Each configuration is trained for the same number of optimizer steps.

This setting matches the one most often used in practice. Unless noted otherwise, hyperparameters are tuned on a held-out validation split. Variance across seeds is small compared with the differences we report. Prior work reached similar conclusions on smaller models.

Results for individual tasks appear in the appendix. We release the evaluation scripts so that every number can be reproduced. We discuss the limitations of this protocol in the final section. We keep the tokenizer and context length fixed across every configuration.

The comparison holds the data and the evaluation harness constant. This setting matches the one most often used in practice. Unless noted otherwise, hyperparameters are tuned on a held-out validation split. Variance across seeds is small compared with the differences we report.

Ablations in the appendix vary one factor at a time. Results for individual tasks appear in the appendix. We release the evaluation scripts so that every number can be reproduced. We discuss the limitations of this protocol in the final section.

We evaluate the final checkpoint rather than the best one seen during training. We report wall-clock time on a single node of eight accelerators. We follow the standard training recipe and report the mean over three seeds. Each configuration is trained for the same number of optimizer steps.

This setting matches the one most often used in practice. Unless noted otherwise, hyperparameters are tuned on a held-out validation split. Variance across seeds is small compared with the differences we report. Prior work reached similar conclusions on smaller models.

Results for individual tasks appear in the appendix. We release the evaluation scripts so that every number can be reproduced. We discuss the limitations of this protocol in the final section. We keep the tokenizer and context length fixed across every configuration.

We report wall-clock time on a single node of eight accelerators. We follow the standard training recipe and report the mean over three seeds. Each configuration is trained for the same number of optimizer steps. Training uses mixed precision with gradient clipping at a norm of one.

Ablations in the appendix vary one factor at a time. Results for individual tasks appear in the appendix. We release the evaluation scripts so that every number can be reproduced. We discuss the limitations of this protocol in the final section.

We evaluate the final checkpoint rather than the best one seen during training. We report wall-clock time on a single node of eight accelerators. We follow the standard training recipe and report the mean over three seeds. Each configuration is trained for the same number of optimizer steps.

This setting matches the one most often used in practice. Unless noted otherwise, hyperparameters are tuned on a held-out validation split. Variance across seeds is small compared with the differences we report. Prior work reached similar conclusions on smaller models.

Results for individual tasks appear in the appendix. We release the evaluation scripts so that every number can be reproduced. We discuss the limitations of this protocol in the final section. We keep the tokenizer and context length fixed across every configuration.

We report wall-clock time on a single node of eight accelerators. We follow the standard training recipe and report the mean over three seeds. Each configuration is trained for the same number of optimizer steps. Training uses mixed precision with gradient clipping at a norm of one.

Unless noted otherwise, hyperparameters are tuned on a held-out validation split. Variance across seeds is small compared with the differences we report. Prior work reached similar conclusions on smaller models. All runs share the same learning-rate schedule with a short linear warmup.

We evaluate the final checkpoint rather than the best one seen during training. We report wall-clock time on a single node of eight accelerators. We follow the standard training recipe and report the mean over three seeds. Each configuration is trained for the same number of optimizer steps.

This setting matches the one most often used in practice. Unless noted otherwise, hyperparameters are tuned on a held-out validation split. Variance across seeds is small compared with the differences we report. Prior work reached similar conclusions on smaller models.

Results for individual tasks appear in the appendix. We release the evaluation scripts so that every number can be reproduced. We discuss the limitations of this protocol in the final section. We keep the tokenizer and context length fixed across every configuration.

We report wall-clock time on a single node of eight accelerators. We follow the standard training recipe and report the mean over three seeds. Each configuration is trained for the same number of optimizer steps. Training uses mixed precision with gradient clipping at a norm of one.

Unless noted otherwise, hyperparameters are tuned on a held-out validation split. Variance across seeds is small compared with the differences we report. Prior work reached similar conclusions on smaller models. All runs share the same learning-rate schedule with a short linear warmup.

We release the evaluation scripts so that every number can be reproduced. We discuss the limitations of this protocol in the final section. We keep the tokenizer and context length fixed across every configuration. The comparison holds the data and the evaluation harness constant.

This setting matches the one most often used in practice. Unless noted otherwise, hyperparameters are tuned on a held-out validation split. Variance across seeds is small compared with the differences we report. Prior work reached similar conclusions on smaller models.

Results for individual tasks appear in the appendix. We release the evaluation scripts so that every number can be reproduced. We discuss the limitations of this protocol in the final section. We keep the tokenizer and context length fixed across every configuration.

We report wall-clock time on a single node of eight accelerators. We follow the standard training recipe and report the mean over three seeds. Each configuration is trained for the same number of optimizer steps. Training uses mixed precision with gradient clipping at a norm of one.

Unless noted otherwise, hyperparameters are tuned on a held-out validation split. Variance across seeds is small compared with the differences we report. Prior work reached similar conclusions on smaller models. All runs share the same learning-rate schedule with a short linear warmup.

We release the evaluation scripts so that every number can be reproduced. We discuss the limitations of this protocol in the final section. We keep the tokenizer and context length fixed across every configuration. The comparison holds the data and the evaluation harness constant.

We follow the standard training recipe and report the mean over three seeds. Each configuration is trained for the same number of optimizer steps. Training uses mixed precision with gradient clipping at a norm of one. Ablations in the appendix vary one factor at a time.

Results for individual tasks appear in the appendix. We release the evaluation scripts so that every number can be reproduced. We discuss the limitations of this protocol in the final section. We keep the tokenizer and context length fixed across every configuration.

We report wall-clock time on a single node of eight accelerators. We follow the standard training recipe and report the mean over three seeds. Each configuration is trained for the same number of optimizer steps. Training uses mixed precision with gradient clipping at a norm of one.

Unless noted otherwise, hyperparameters are tuned on a held-out validation split. Variance across seeds is small compared with the differences we report. Prior work reached similar conclusions on smaller models. All runs share the same learning-rate schedule with a short linear warmup.

We release the evaluation scripts so that every number can be reproduced. We discuss the limitations of this protocol in the final section. We keep the tokenizer and context length fixed across every configuration. The comparison holds the data and the evaluation harness constant.

We follow the standard training recipe and report the mean over three seeds. Each configuration is trained for the same number of optimizer steps. Training uses mixed precision with gradient clipping at a norm of one. Ablations in the appendix vary one factor at a time.

Variance across seeds is small compared with the differences we report. Prior work reached similar conclusions on smaller models. All runs share the same learning-rate schedule with a short linear warmup. We evaluate the final checkpoint rather than the best one seen during training.

We report wall-clock time on a single node of eight accelerators. We follow the standard training recipe and report the mean over three seeds. Each configuration is trained for the same number of optimizer steps. Training uses mixed precision with gradient clipping at a norm of one.

Unless noted otherwise, hyperparameters are tuned on a held-out validation split. Variance across seeds is small compared with the differences we report. Prior work reached similar conclusions on smaller models. All runs share the same learning-rate schedule with a short linear warmup.

We release the evaluation scripts so that every number can be reproduced. We discuss the limitations of this protocol in the final section. We keep the tokenizer and context length fixed across every configuration. The comparison holds the data and the evaluation harness constant.

We follow the standard training recipe and report the mean over three seeds. Each configuration is trained for the same number of optimizer steps. Training uses mixed precision with gradient clipping at a norm of one. Ablations in the appendix vary one factor at a time.

Variance across seeds is small compared with the differences we report. Prior work reached similar conclusions on smaller models. All runs share the same learning-rate schedule with a short linear warmup. We evaluate the final checkpoint rather than the best one seen during training.

We discuss the limitations of this protocol in the final section. We keep the tokenizer and context length fixed across every configuration. The comparison holds the data and the evaluation harness constant. This setting matches the one most often used in practice.

Unless noted otherwise, hyperparameters are tuned on a held-out validation split. Variance across seeds is small compared with the differences we report. Prior work reached similar conclusions on smaller models. All runs share the same learning-rate schedule with a short linear warmup.

We release the evaluation scripts so that every number can be reproduced. We discuss the limitations of this protocol in the final section. We keep the tokenizer and context length fixed across every configuration. The comparison holds the data and the evaluation harness constant.

We follow the standard training recipe and report the mean over three seeds. Each configuration is trained for the same number of optimizer steps. Training uses mixed precision with gradient clipping at a norm of one. Ablations in the appendix vary one factor at a time.

Variance across seeds is small compared with the differences we report. Prior work reached similar conclusions on smaller models. All runs share the same learning-rate schedule with a short linear warmup. We evaluate the final checkpoint rather than the best one seen during training.

We discuss the limitations of this protocol in the final section. We keep the tokenizer and context length fixed across every configuration. The comparison holds the data and the evaluation harness constant. This setting matches the one most often used in practice.

Each configuration is trained for the same number of optimizer steps. Training uses mixed precision with gradient clipping at a norm of one. Ablations in the appendix vary one factor at a time. Results for individual tasks appear in the appendix.

We release the evaluation scripts so that every number can be reproduced. We discuss the limitations of this protocol in the final section. We keep the tokenizer and context length fixed across every configuration. The comparison holds the data and the evaluation harness constant.

We follow the standard training recipe and report the mean over three seeds. Each configuration is trained for the same number of optimizer steps. Training uses mixed precision with gradient clipping at a norm of one. Ablations in the appendix vary one factor at a time.

Variance across seeds is small compared with the differences we report. Prior work reached similar conclusions on smaller models. All runs share the same learning-rate schedule with a short linear warmup. We evaluate the final checkpoint rather than the best one seen during training.

We discuss the limitations of this protocol in the final section. We keep the tokenizer and context length fixed across every configuration. The comparison holds the data and the evaluation harness constant. This setting matches the one most often used in practice.

Each configuration is trained for the same number of optimizer steps. Training uses mixed precision with gradient clipping at a norm of one. Ablations in the appendix vary one factor at a time. Results for individual tasks appear in the appendix.

Prior work reached similar conclusions on smaller models. All runs share the same learning-rate schedule with a short linear warmup. We evaluate the final checkpoint rather than the best one seen during training. We report wall-clock time on a single node of eight accelerators.

We follow the standard training recipe and report the mean over three seeds. Each configuration is trained for the same number of optimizer steps. Training uses mixed precision with gradient clipping at a norm of one. Ablations in the appendix vary one factor at a time.

Variance across seeds is small compared with the differences we report. Prior work reached similar conclusions on smaller models. All runs share the same learning-rate schedule with a short linear warmup. We evaluate the final checkpoint rather than the best one seen during training.

We discuss the limitations of this protocol in the final section. We keep the tokenizer and context length fixed across every configuration. The comparison holds the data and the evaluation harness constant. This setting matches the one most often used in practice.

Each configuration is trained for the same number of optimizer steps. Training uses mixed precision with gradient clipping at a norm of one. Ablations in the appendix vary one factor at a time. Results for individual tasks appear in the appendix.

Prior work reached similar conclusions on smaller models. All runs share the same learning-rate schedule with a short linear warmup. We evaluate the final checkpoint rather than the best one seen during training. We report wall-clock time on a single node of eight accelerators.

We keep the tokenizer and context length fixed across every configuration. The comparison holds the data and the evaluation harness constant. This setting matches the one most often used in practice. Unless noted otherwise, hyperparameters are tuned on a held-out validation split.

Variance across seeds is small compared with the differences we report. Prior work reached similar conclusions on smaller models. All runs share the same learning-rate schedule with a short linear warmup. We evaluate the final checkpoint rather than the best one seen during training.

We discuss the limitations of this protocol in the final section. We keep the tokenizer and context length fixed across every configuration. The comparison holds the data and the evaluation harness constant. This setting matches the one most often used in practice.

Each configuration is trained for the same number of optimizer steps. Training uses mixed precision with gradient clipping at a norm of one. Ablations in the appendix vary one factor at a time. Results for individual tasks appear in the appendix.

Prior work reached similar conclusions on smaller models. All runs share the same learning-rate schedule with a short linear warmup. We evaluate the final checkpoint rather than the best one seen during training. We report wall-clock time on a single node of eight accelerators.

We keep the tokenizer and context length fixed across every configuration. The comparison holds the data and the evaluation harness constant. This setting matches the one most often used in practice. Unless noted otherwise, hyperparameters are tuned on a held-out validation split.

Training uses mixed precision with gradient clipping at a norm of one. Ablations in the appendix vary one factor at a time. Results for individual tasks appear in the appendix. We release the evaluation scripts so that every number can be reproduced.

We discuss the limitations of this protocol in the final section. We keep the tokenizer and context length fixed across every configuration. The comparison holds the data and the evaluation harness constant. This setting matches the one most often used in practice.

Each configuration is trained for the same number of optimizer steps. Training uses mixed precision with gradient clipping at a norm of one. Ablations in the appendix vary one factor at a time. Results for individual tasks appear in the appendix.

Prior work reached similar conclusions on smaller models. All runs share the same learning-rate schedule with a short linear warmup. We evaluate the final checkpoint rather than the best one seen during training. We report wall-clock time on a single node of eight accelerators.

We keep the tokenizer and context length fixed across every configuration. The comparison holds the data and the evaluation harness constant. This setting matches the one most often used in practice. Unless noted otherwise, hyperparameters are tuned on a held-out validation split.

Training uses mixed precision with gradient clipping at a norm of one. Ablations in the appendix vary one factor at a time. Results for individual tasks appear in the appendix. We release the evaluation scripts so that every number can be reproduced.

All runs share the same learning-rate schedule with a short linear warmup. We evaluate the final checkpoint rather than the best one seen during training. We report wall-clock time on a single node of eight accelerators. We follow the standard training recipe and report the mean over three seeds.

Each configuration is trained for the same number of optimizer steps. Training uses mixed precision with gradient clipping at a norm of one. Ablations in the appendix vary one factor at a time. Results for individual tasks appear in the appendix.

Prior work reached similar conclusions on smaller models. All runs share the same learning-rate schedule with a short linear warmup. We evaluate the final checkpoint rather than the best one seen during training. We report wall-clock time on a single node of eight accelerators.

We keep the tokenizer and context length fixed across every configuration. The comparison holds the data and the evaluation harness constant. This setting matches the one most often used in practice. Unless noted otherwise, hyperparameters are tuned on a held-out validation split.

Training uses mixed precision with gradient clipping at a norm of one. Ablations in the appendix vary one factor at a time. Results for individual tasks appear in the appendix. We release the evaluation scripts so that every number can be reproduced.

All runs share the same learning-rate schedule with a short linear warmup. We evaluate the final checkpoint rather than the best one seen during training. We report wall-clock time on a single node of eight accelerators. We follow the standard training recipe and report the mean over three seeds.

The comparison holds the data and the evaluation harness constant. This setting matches the one most often used in practice. Unless noted otherwise, hyperparameters are tuned on a held-out validation split. Variance across seeds is small compared with the differences we report.

Prior work reached similar conclusions on smaller models. All runs share the same learning-rate schedule with a short linear warmup. We evaluate the final checkpoint rather than the best one seen during training. We report wall-clock time on a single node of eight accelerators.

We keep the tokenizer and context length fixed across every configuration. The comparison holds the data and the evaluation harness constant. This setting matches the one most often used in practice. Unless noted otherwise, hyperparameters are tuned on a held-out validation split.

Training uses mixed precision with gradient clipping at a norm of one. Ablations in the appendix vary one factor at a time. Results for individual tasks appear in the appendix. We release the evaluation scripts so that every number can be reproduced.

All runs share the same learning-rate schedule with a short linear warmup. We evaluate the final checkpoint rather than the best one seen during training. We report wall-clock time on a single node of eight accelerators. We follow the standard training recipe and report the mean over three seeds.

The comparison holds the data and the evaluation harness constant. This setting matches the one most often used in practice. Unless noted otherwise, hyperparameters are tuned on a held-out validation split. Variance across seeds is small compared with the differences we report.

Ablations in the appendix vary one factor at a time. Results for individual tasks appear in the appendix. We release the evaluation scripts so that every number can be reproduced. We discuss the limitations of this protocol in the final section.

We keep the tokenizer and context length fixed across every configuration. The comparison holds the data and the evaluation harness constant. This setting matches the one most often used in practice. Unless noted otherwise, hyperparameters are tuned on a held-out validation split.

Training uses mixed precision with gradient clipping at a norm of one. Ablations in the appendix vary one factor at a time. Results for individual tasks appear in the appendix. We release the evaluation scripts so that every number can be reproduced.

All runs share the same learning-rate schedule with a short linear warmup. We evaluate the final checkpoint rather than the best one seen during training. We report wall-clock time on a single node of eight accelerators. We follow the standard training recipe and report the mean over three seeds.

The comparison holds the data and the evaluation harness constant. This setting matches the one most often used in practice. Unless noted otherwise, hyperparameters are tuned on a held-out validation split. Variance across seeds is small compared with the differences we report.

Ablations in the appendix vary one factor at a time. Results for individual tasks appear in the appendix. We release the evaluation scripts so that every number can be reproduced. We discuss the limitations of this protocol in the final section.

We evaluate the final checkpoint rather than the best one seen during training. We report wall-clock time on a single node of eight accelerators. We follow the standard training recipe and report the mean over three seeds. Each configuration is trained for the same number of optimizer steps.

Training uses mixed precision with gradient clipping at a norm of one. Ablations in the appendix vary one factor at a time. Results for individual tasks appear in the appendix. We release the evaluation scripts so that every number can be reproduced.

All runs share the same learning-rate schedule with a short linear warmup. We evaluate the final checkpoint rather than the best one seen during training. We report wall-clock time on a single node of eight accelerators. We follow the standard training recipe and report the mean over three seeds.

The comparison holds the data and the evaluation harness constant. This setting matches the one most often used in practice. Unless noted otherwise, hyperparameters are tuned on a held-out validation split. Variance across seeds is small compared with the differences we report.

Ablations in the appendix vary one factor at a time. Results for individual tasks appear in the appendix. We release the evaluation scripts so that every number can be reproduced. We discuss the limitations of this protocol in the final section.

We evaluate the final checkpoint rather than the best one seen during training. We report wall-clock time on a single node of eight accelerators. We follow the standard training recipe and report the mean over three seeds. Each configuration is trained for the same number of optimizer steps.

This setting matches the one most often used in practice. Unless noted otherwise, hyperparameters are tuned on a held-out validation split. Variance across seeds is small compared with the differences we report. Prior work reached similar conclusions on smaller models.

All runs share the same learning-rate schedule with a short linear warmup. We evaluate the final checkpoint rather than the best one seen during training. We report wall-clock time on a single node of eight accelerators. We follow the standard training recipe and report the mean over three seeds.

The comparison holds the data and the evaluation harness constant. This setting matches the one most often used in practice. Unless noted otherwise, hyperparameters are tuned on a held-out validation split. Variance across seeds is small compared with the differences we report.

Ablations in the appendix vary one factor at a time. Results for individual tasks appear in the appendix. We release the evaluation scripts so that every number can be reproduced. We discuss the limitations of this protocol in the final section.

We evaluate the final checkpoint rather than the best one seen during training. We report wall-clock time on a single node of eight accelerators. We follow the standard training recipe and report the mean over three seeds. Each configuration is trained for the same number of optimizer steps.

This setting matches the one most often used in practice. Unless noted otherwise, hyperparameters are tuned on a held-out validation split. Variance across seeds is small compared with the differences we report. Prior work reached similar conclusions on smaller models.

Results for individual tasks appear in the appendix. We release the evaluation scripts so that every number can be reproduced. We discuss the limitations of this protocol in the final section. We keep the tokenizer and context length fixed across every configuration.

The comparison holds the data and the evaluation harness constant. This setting matches the one most often used in practice. Unless noted otherwise, hyperparameters are tuned on a held-out validation split. Variance across seeds is small compared with the differences we report.

Ablations in the appendix vary one factor at a time. Results for individual tasks appear in the appendix. We release the evaluation scripts so that every number can be reproduced. We discuss the limitations of this protocol in the final section.

We evaluate the final checkpoint rather than the best one seen during training. We report wall-clock time on a single node of eight accelerators. We follow the standard training recipe and report the mean over three seeds. Each configuration is trained for the same number of optimizer steps.

This setting matches the one most often used in practice. Unless noted otherwise, hyperparameters are tuned on a held-out validation split. Variance across seeds is small compared with the differences we report. Prior work reached similar conclusions on smaller models.

Results for individual tasks appear in the appendix. We release the evaluation scripts so that every number can be reproduced. We discuss the limitations of this protocol in the final section. We keep the tokenizer and context length fixed across every configuration.

We report wall-clock time on a single node of eight accelerators. We follow the standard training recipe and report the mean over three seeds. Each configuration is trained for the same number of optimizer steps. Training uses mixed precision with gradient clipping at a norm of one.

Ablations in the appendix vary one factor at a time. Results for individual tasks appear in the appendix. We release the evaluation scripts so that every number can be reproduced. We discuss the limitations of this protocol in the final section.

We evaluate the final checkpoint rather than the best one seen during training. We report wall-clock time on a single node of eight accelerators. We follow the standard training recipe and report the mean over three seeds. Each configuration is trained for the same number of optimizer steps.

This setting matches the one most often used in practice. Unless noted otherwise, hyperparameters are tuned on a held-out validation split. Variance across seeds is small compared with the differences we report. Prior work reached similar conclusions on smaller models.

Results for individual tasks appear in the appendix. We release the evaluation scripts so that every number can be reproduced. We discuss the limitations of this protocol in the final section. We keep the tokenizer and context length fixed across every configuration.

We report wall-clock time on a single node of eight accelerators. We follow the standard training recipe and report the mean over three seeds. Each configuration is trained for the same number of optimizer steps. Training uses mixed precision with gradient clipping at a norm of one.

Unless noted otherwise, hyperparameters are tuned on a held-out validation split. Variance across seeds is small compared with the differences we report. Prior work reached similar conclusions on smaller models. All runs share the same learning-rate schedule with a short linear warmup.

We evaluate the final checkpoint rather than the best one seen during training. We report wall-clock time on a single node of eight accelerators. We follow the standard training recipe and report the mean over three seeds. Each configuration is trained for the same number of optimizer steps.

This setting matches the one most often used in practice. Unless noted otherwise, hyperparameters are tuned on a held-out validation split. Variance across seeds is small compared with the differences we report. Prior work reached similar conclusions on smaller models.

Results for individual tasks appear in the appendix. We release the evaluation scripts so that every number can be reproduced. We discuss the limitations of this protocol in the final section. We keep the tokenizer and context length fixed across every configuration.

We report wall-clock time on a single node of eight accelerators. We follow the standard training recipe and report the mean over three seeds. Each configuration is trained for the same number of optimizer steps. Training uses mixed precision with gradient clipping at a norm of one.

Unless noted otherwise, hyperparameters are tuned on a held-out validation split. Variance across seeds is small compared with the differences we report. Prior work reached similar conclusions on smaller models. All runs share the same learning-rate schedule with a short linear warmup.

We release the evaluation scripts so that every number can be reproduced. We discuss the limitations of this protocol in the final section. We keep the tokenizer and context length fixed across every configuration. The comparison holds the data and the evaluation harness constant.

This setting matches the one most often used in practice. Unless noted otherwise, hyperparameters are tuned on a held-out validation split. Variance across seeds is small compared with the differences we report. Prior work reached similar conclusions on smaller models.

Results for individual tasks appear in the appendix. We release the evaluation scripts so that every number can be reproduced. We discuss the limitations of this protocol in the final section. We keep the tokenizer and context length fixed across every configuration.

We report wall-clock time on a single node of eight accelerators. We follow the standard training recipe and report the mean over three seeds. Each configuration is trained for the same number of optimizer steps. Training uses mixed precision with gradient clipping at a norm of one.

Unless noted otherwise, hyperparameters are tuned on a held-out validation split. Variance across seeds is small compared with the differences we report. Prior work reached similar conclusions on smaller models. All runs share the same learning-rate schedule with a short linear warmup.

We release the evaluation scripts so that every number can be reproduced. We discuss the limitations of this protocol in the final section. We keep the tokenizer and context length fixed across every configuration. The comparison holds the data and the evaluation harness constant.

We follow the standard training recipe and report the mean over three seeds. Each configuration is trained for the same number of optimizer steps. Training uses mixed precision with gradient clipping at a norm of one. Ablations in the appendix vary one factor at a time.

Results for individual tasks appear in the appendix. We release the evaluation scripts so that every number can be reproduced. We discuss the limitations of this protocol in the final section. We keep the tokenizer and context length fixed across every configuration.

We report wall-clock time on a single node of eight accelerators. We follow the standard training recipe and report the mean over three seeds. Each configuration is trained for the same number of optimizer steps. Training uses mixed precision with gradient clipping at a norm of one.

Unless noted otherwise, hyperparameters are tuned on a held-out validation split. Variance across seeds is small compared with the differences we report. Prior work reached similar conclusions on smaller models. All runs share the same learning-rate schedule with a short linear warmup.

We release the evaluation scripts so that every number can be reproduced. We discuss the limitations of this protocol in the final section. We keep the tokenizer and context length fixed across every configuration. The comparison holds the data and the evaluation harness constant.

We follow the standard training recipe and report the mean over three seeds. Each configuration is trained for the same number of optimizer steps. Training uses mixed precision with gradient clipping at a norm of one. Ablations in the appendix vary one factor at a time.

Variance across seeds is small compared with the differences we report. Prior work reached similar conclusions on smaller models. All runs share the same learning-rate schedule with a short linear warmup. We evaluate the final checkpoint rather than the best one seen during training.